Intro to LLM's

Monday, March 16, 2026 8:23 PM

Course Overview

The course targets developers using LLM APIs to build applications - not just the ChatGPT web UI. Topics covered:

- Prompting best practices for software development
- Common use cases: summarizing, inferring, transforming, expanding
- Building a chatbot with an LLM

Two Types of LLMs

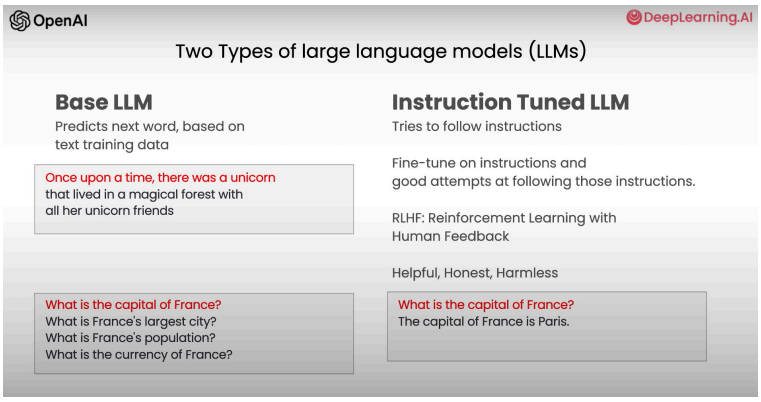

Base LLMs

- Trained to predict the next word from large internet/text datasets
- Example: prompt "Once upon a time there was a unicorn..." → continues the story
- Can behave unexpectedly with direct questions (e.g., "What is the capital of France?" might return a list of quiz questions)

Instruction-Tuned LLMs *(recommended for most applications)*

- Trained to follow instructions directly
- Same question → "The capital of France is Paris"
- Training process: start with a base LLM → fine-tune on instruction/response pairs → refine with **RLHF**(Reinforcement Learning from Human Feedback)
- Goals: helpful, honest, and harmless; less likely to produce toxic outputs

Key Prompting Principles (previewed)

1. **Be clear and specific** - vague prompts yield vague results
2. **Give the LLM time to think** - covered in the next lesson

Practical tip: Think of prompting like giving instructions to a smart person who doesn't know your specific task. The more context you provide (tone, focus, background material), the better the output.