Expanding

Monday, March 16, 2026 11:41 PM

What is Expanding?

Taking a short input (instructions, bullet points, a topic) and having the LLM generate a longer piece of text like an email or essay. Great for brainstorming and drafting, but can be misused (e.g. spam generation), so use responsibly.

Practical Example: Personalized Customer Email

Given a customer review and its sentiment (already extracted from the inferring step), generate a custom reply email.

Prompt instructions used:

- Role: "You are a customer service AI assistant"
- If sentiment is positive/neutral: thank the customer
- If sentiment is negative: apologize and direct them to customer service
- Use specific details from the review
- Keep a concise, professional tone
- Sign off as "AI customer agent" (transparency that the reply is AI-generated)
  Always disclose when text shown to a user was generated by AI. This is an ethical responsibility, not just a best practice.

Temperature

Temperature controls how random or varied the model's outputs are. Think of it as a creativity dial.

How it works conceptually:

When predicting the next word, the model assigns probabilities to all possible words. Temperature determines how strictly it follows those probabilities.

Temperature	Behavior	Best For
0	Always picks the highest probability word. Same prompt gives same output every time.	Reliable, predictable apps
0.7	Occasionally picks lower probability words. Output varies between runs.	Drafting, brainstorming
High (e.g. 1.0 +)	Can pick unlikely words. More creative but less consistent.	Creative writing, ideation

A useful mental model: low temperature = focused and consistent, high temperature = distracted but creative. As the model generates more tokens, small early differences compound into very different final outputs.

Recommendation: Use temperature 0 for production systems where reliability matters. Use higher temperatures when you want variety or creative output.

Key Takeaway

Expanding is powerful for automating drafting tasks. Pair it with outputs from earlier steps (like extracted sentiment) to build multi-step pipelines. Always tune temperature based on whether you need consistency or creativity.

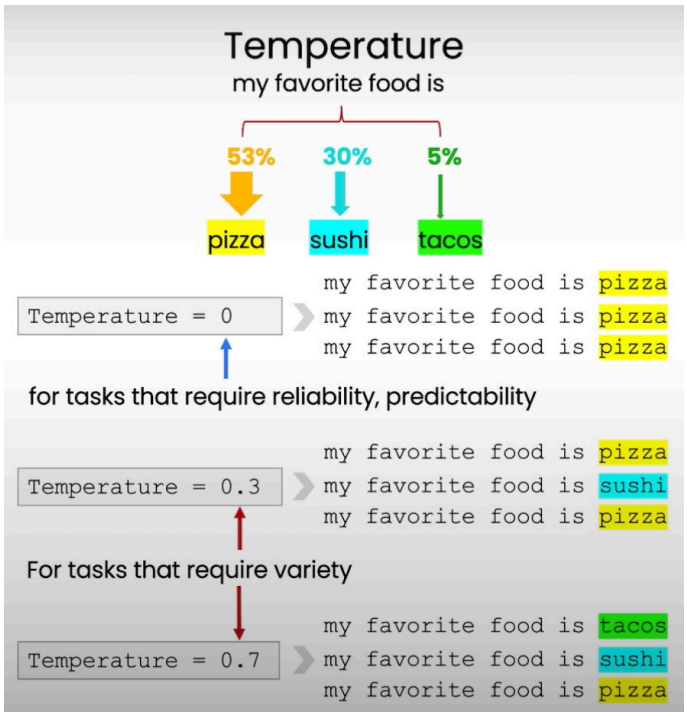
Temperature
my favorite food is
53%
30%
5%
pizza
sushi
tacos
Temperature = 0
my favorite food is pizza
my favorite food is pizza
my favorite food is pizza
for tasks that require reliability, predictability
Temperature = 0.3
my favorite food is pizza
my favorite food is sushi
my favorite food is pizza
For tasks that require variety
Temperature = 0.7
my favorite food is tacos
my favorite food is sushi
my favorite food is pizza